

Kabeer Singh Thockchom

+1-530-564-9326 | thockchomkabeer@gmail.com |

 [kabeerthockchom](#) |  [KabeerThockchom](#) |  [Kabeerthockchom.dev](#) |

San Francisco, CA

SUMMARY

Senior Solutions Engineer with **4+ years** of customer-facing AI experience, currently driving **greenfield account strategy** across **21+ Fortune 500 retail, travel, and hospitality accounts** on the **Databricks Lakehouse**. Run technical discovery with C-suite buyers, architect end-to-end migrations off Snowflake / BigQuery / Domo / Power BI / AS400, and ship customer-facing demos and reusable reference assets that move accounts from pre-discovery to committed POC in weeks. Hands-on builder of **multi-agent platforms, LLM-powered demos, and RAG architectures** that have shipped to **10+ Fortune 500 customers**. Background combines **AI platform engineering** (semantic search, MLOps, evaluation frameworks at \$10T AUM clients), **startup founding** (5 F500 deployments), and **stage / executive presence** (NVIDIA GTC 2025 speaker, hackathon winner).

EXPERIENCE

• Databricks

December 2025 - Present

Senior Solutions Engineer – Retail, Travel & Hospitality

San Francisco, CA

- **Greenfield Account Ownership & Consumption Growth:** Inherited **21+ greenfield F500 retail / travel / hospitality / healthcare / food-distribution accounts** across **two AEs**, starting from \$0 ARR and zero workspace footprint → own technical strategy, drive consumption, and convert pre-discovery into committed POCs → ran territory planning with AEs, qualified opportunities, designed migration architectures off Snowflake / BigQuery / Domo / Power BI / **AS400**, scoped CDC ingestion patterns, vector search, and agent deployment → **multiple multi-month POCs in flight** with locked migration plans and committed evaluation criteria
- **Executive Discovery & On-Site Workshops:** F500 membership warehouse retailer running Domo and Power BI across **5,000+ dashboard consumers**, evaluating consolidation onto a single AI platform → earn POC commit and define mutual close plan → led multi-day **on-site discovery workshop in Latin America** with C-suite and VP-level technical leaders, mapped current architecture, ran value framing for AI/BI + Genie + Agent Bricks, and locked the POC sequence (SA workshop → data domain + arch discovery → on-site workshop) → committed POC with phased success criteria; pattern reused across other F500 retail accounts
- **Flagship Demo Engineering – Smallville Mart:** Retail customers needed to see a **believable agentic operations story** grounded in real streaming data, not slideware → build a flagship reusable demo SAs across the org could re-skin per customer → designed Stanford Smallville-inspired retail simulation with **LLM customer agents** making autonomous purchasing decisions, streaming events into a **multi-agent ops platform** (forecasting, dynamic pricing, inventory, labor) with **human-in-the-loop** approvals; authored **5 technical PRDs** covering simulation engine, supervisor / sub-agent architecture, and bidirectional data loops → flagship retail discovery asset used to advance executive conversations and accelerate POC commits
- **Customer-Facing Solution Demos:** F500 retail and hospitality customers questioning whether the Databricks Agent Framework was production-ready → ship reference solutions tailored to live discovery conversations → built a **Markdown Optimization agent** with ReAct reasoning for apparel pricing decisions, a **restaurant inventory ReAct agent** with autonomous tool-use, a **Genie-powered text-to-SQL analytics layer**, and a **browser AI agents reference deployment (CoDA)** on Databricks Apps with full Make-based deployment automation → demos directly contributed to **multiple POCs moving to commit phase** across F500 retail / hospitality accounts
- **Multi-Agent Reference Architecture:** F500 retailers running siloed forecasting, pricing, inventory, and labor systems with no shared intelligence layer → design a reusable Lakehouse-native blueprint other SAs could deliver → architected **Supervisor Agent pattern** with specialized sub-agents and end-to-end pipelines (ingestion → **medallion architecture** → **Unity Catalog** governance → Model Serving → agent execution), grounded in customer data via Genie and UC Functions → reference architecture reused across multiple retail POCs, cutting customer evaluation cycles from months to weeks

- **Internal Enablement & SA Productivity Asset:** Solutions Engineering org adopting Claude Code + MCP without a shared workflow, with each SA rebuilding the same scaffolding → ship reusable internal assets adopted across the SE org → authored team-wide Claude Code + MCP setup guide, automated workspace provisioning via **Govibe plugins**, and built an Obsidian-integrated SA knowledge system covering **20+ account dossiers** with **30+ custom skills** for territory ops, account research, deck building, and data exploration → adopted across SE peers as best-practice template for AI-native field work

• Ernst and Young LLP

July 2025 - December 2025

Senior Applied AI Engineer

San Francisco, CA

- **AI Platform Architecture:** Led enterprise AI platform strategy across **15 engineering teams**; designed **MLOps infrastructure** (model registry, experiment tracking, feature stores, CI/CD pipelines) increasing AI project success rate from 20% to **65%** and reducing model deployment time by **75%**
- **Enterprise RAG Search – \$10T AUM Client:** Architected retrieval pipeline with **semantic search** (cosine similarity, TF-IDF, metadata filtering) using **Qdrant** vector database across 200+ knowledge systems; achieved **90% query accuracy**, reduced resolution from 45 min to **5 min**, processing **10K+ daily queries** with model distillation
- **EY Voice – AI Call Center Platform:** Engineered real-time **speech-to-text pipeline**, automated ticket classification using LLMs, and built agent-assist system; reduced call resolution **45%**, automated **75%** of tickets; presented at **NVIDIA GTC 2025**; secured **\$1.2M funding**
- **AI Agent Evaluation Framework:** Defined **LLM performance benchmarks**, built automated testing pipelines for agent reliability, established governance policies; increased agent efficiency **80%**, reduced onboarding from 2 weeks to **24 hours**
- **NLP Infrastructure:** Built **NER-based** compliance extraction system reducing manpower **90%**; generated **200K+ synthetic datasets** for model fine-tuning and evaluation

• Ernst and Young LLP

October 2023 - July 2025

AI Platform Engineer

San Francisco, CA

- **Semantic Search Infrastructure:** Architected search system using **Qdrant** vector database with cosine similarity scoring, TF-IDF ranking, and metadata filtering; designed **embedding pipelines** for document chunking, indexing, and retrieval across **200+ enterprise systems**
- **AI Agent Framework:** Developed agent framework with **tool-use capabilities** for banking clients; implemented **ReAct-style reasoning loops**, built evaluation harnesses for LLM output quality, created onboarding automation reducing agent setup from 2 weeks to **24 hours**
- **Platform Telemetry & Analytics:** Built dashboards tracking model performance, API latency, retrieval accuracy, and agent success rates across **200+ hours** of system analysis; drove platform **NPS from 20 to 75**; secured **\$3M additional investment**
- **Engineering Leadership:** Managed **15 engineers** across **3 locations**, optimized ML pipeline velocity by **40%** using JIRA/Confluence following SAFe Agile methodology

• Stealth Startup

January 2024 - July 2025

Founder & AI Engineer

San Francisco, CA

- **EYLAR – Enterprise RAG Platform:** Built on-premise RAG platform with **zero-trust architecture**: document parsing (PDF, PPT, Word, Excel), chunking strategies, embedding generation, and **vector retrieval** with transparent citations; **60% latency reduction**, **35% accuracy improvement**, deployed across **5 Fortune 500** companies
- **Bottega AI – Restaurant AI Agent:** Developed AI voice agent using **ReAct architecture** with tool-use for menu querying, order processing, and upselling; achieved **40%** reduction in order processing, **25%** upselling increase through personalization
- **Talk2Data – Text-to-SQL Analytics:** Architected natural language **Text-to-SQL** platform for ETF firm: query parsing, schema inference, SQL generation, and result visualization processing **10K+ daily queries** with **95% accuracy**

• Ernst and Young LLP

June 2022 - August 2022

AI Engineering Intern

San Francisco, CA

- Built data pipelines and process workflows for **CECL** credit risk framework with **30+ stakeholders**; developed **6 process workflows** with risk controls reducing implementation by **30%**; achieved **100% compliance**

• UC Davis Graduate School of Management

June 2021 - August 2022

Research Analyst – NLP

Davis, CA

- Led **10-person team** developing NLP classification model using **BERT fine-tuning**; designed annotation guidelines, built labeled datasets, ran A/B testing improving accuracy **15%** and labeling consistency **25%**

CUSTOMER-FACING DEMOS & REUSABLE ASSETS

• Smallville Mart – Generative Agents Retail Simulation

Feb 2026 - Present

Impact: flagship F500 retail discovery demo, reused across SE org

[Overview](#)

- Stanford Smallville-inspired simulation with **LLM customer agents** streaming live events into a multi-agent ops platform (forecasting, pricing, inventory, labor) with **human-in-the-loop** approvals; 5 PRDs for engine, agents, and data loops

• Multi-Agent Retail Ops Reference Architecture

Jan 2026 - Present

Impact: reusable blueprint accelerating retail POCs from months to weeks

Internal asset

- **Supervisor Agent pattern** with specialized sub-agents on Databricks Lakehouse: ingestion → medallion → Unity Catalog → Model Serving → agent execution; reused across multiple F500 retail POCs

• Browser AI Agents on Databricks (CoDA)

Mar 2026 - Present

Impact: live customer-facing demo for agentic discovery conversations

Internal demo

- Reference deployment of browser-controlling AI agents on **Databricks Apps** with end-to-end Make-based deployment automation on a managed workspace

• Text2SQL – Natural Language Analytics Engine

Jan 2024 - Present

Impact: 500+ users, 70% analyst workload reduction

[Live Demo](#)

- Adaptive **SQL generation system** with 100+ visualization types, query intent classification, schema-aware prompt engineering, continuous learning feedback loops

• EY Voice – Real-Time AI Call Center Platform

Oct 2023 - Present

Impact: 10+ enterprise clients, 1M+ calls/month

[Live Demo](#)

- Speech-to-text + LLM pipeline for real-time call analysis, automated ticket routing, agent-assist; **NVIDIA GTC 2025** keynote; secured **\$1.2M** funding

• EYLAR – Enterprise Agentic RAG Platform

Nov 2023 - Present

Impact: 5 Fortune 500 deployments

[Overview](#)

- End-to-end RAG with zero-trust architecture, multi-format document processing, Qdrant vector search, transparent citations

EDUCATION

• University of Texas, Austin

January 2025 - Present

M.S. in Data Science

Austin, TX

- Focus: Machine Learning Systems, Distributed Computing, Statistical Learning
- Coursework: MLOps, A/B Testing, Experiment Design

• University of California, Davis

September 2019 - June 2023

B.S. Computer Science & Quantitative Economics

Davis, CA

- GPA: **3.7/4.0** | Minors: Statistics & Technology Management
- Honors: Dean's Scholar (2020-2022), March Fund Award (2022)
- Coursework: Machine Learning, Database Systems, Statistical Analysis, HCI

SKILLS

- **Solutions Engineering:** Greenfield account development, executive discovery (CIO/CTO/CDO), POC scoping & delivery, migration architecture (AS400, Snowflake, BigQuery, Domo, Power BI → Databricks), demo / reference asset engineering, on-site workshops, AE partnership, territory planning, MEDDPICCC, value framing
- **Databricks Platform:** Unity Catalog, Model Serving, Lakeflow / DLT, Agent Bricks & Agent Framework, Genie, AI/BI Dashboards, Databricks Apps, Asset Bundles (DABs), Lakebase, MCP, Vector Search
- **LLMs & Agents:** GPT-4, Claude, LLaMA, RAG architectures, ReAct agents, tool-use patterns, multi-agent orchestration (Supervisor pattern), prompt engineering, fine-tuning (LoRA/QLoRA), evaluation frameworks
- **Data & ML:** Python, SQL, Delta Lake, Apache Spark, PyTorch, TensorFlow, Hugging Face Transformers, LangChain, MLflow, Qdrant, Pinecone, FAISS
- **Cloud & Infra:** AWS, Azure, GCP, Docker, Kubernetes, CI/CD, Snowflake, BigQuery, PostgreSQL, MongoDB
- **MLOps:** Model registry, experiment tracking, feature stores, model monitoring, A/B testing, LLMOps
- **Languages (Spoken):** English (Native), Hindi (Native), Spanish (Conversational), Punjabi (Conversational)

CERTIFICATIONS

- | | |
|---|------|
| • Project Management Professional (PMP) | 2024 |
| • Microsoft AI-900 Azure AI Fundamentals | 2024 |
| • DeepLearning.AI LLM Fine-tuning Specialization | 2023 |

RECOGNITION & SPEAKING

- | | |
|--|--|
| • NVIDIA GTC 2025 – Featured Speaker
“EY Voice: AI-Powered Contact Centers” | March 2025
Conference |
| ◦ Presented enterprise AI platform architecture to 500+ attendees , demonstrating 45% call reduction metrics | |
| • EY Wealth Asset Management Hackathon Winner
Ernst and Young LLP | February 2025 |
| ◦ Developed multi-agent platform automating wealth management workflows, won \$50K prize | |
| • Dean’s Scholar & March Fund Award
University of California Davis | 2020-2022 |
| ◦ Top 16% of class for academic excellence and leadership impact | |